

Question 1)

```
# 1. Import Libraries
import pandas as pd
import numpy as np
import seaborn as sn
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

# 2. Creating the 'sales' Dataset
np.random.seed(0) # For reproducibility
data = pd.DataFrame({
    'ID': range(1, 501),
    'TV': np.random.uniform(50, 300, 500),
    'Radio': np.random.uniform(10, 100, 500),
    'Newspaper': np.random.uniform(5, 50, 500),
    'Sales': np.random.uniform(5, 25, 500)
})
print(data.head()) # Previewing the first 5 rows

# 3. Splitting the Data into Independent and Target Variables
X = data[['TV', 'Radio', 'Newspaper']].values # Independent variables
y = data['Sales'].values # Target variable

# Splitting into Training and Testing Sets (7:3 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=0)

print("Training Set Shape:", X_train.shape, y_train.shape)
print("Testing Set Shape:", X_test.shape, y_test.shape)

# 4. Building a Linear Regression Model
linear_regressor = LinearRegression() # Creating the linear
regression model
linear_regressor.fit(X_train, y_train) # Fitting the model

# Making predictions on the test set
y_pred = linear_regressor.predict(X_test)

# 5. Evaluating the Model
print("Model Coefficients:", linear_regressor.coef_)
print("Model Intercept:", linear_regressor.intercept_)

# Calculating R-squared and Mean Squared Error
from sklearn.metrics import r2_score, mean_squared_error
r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)


print(f"R-squared: {r2:.2f}")
```

```
print(f"Mean Squared Error: {mse:.2f}")

# Plotting Actual vs Predicted Sales
plt.scatter(y_test, y_pred, color='blue', alpha=0.6)
plt.xlabel('Actual Sales')
plt.ylabel('Predicted Sales')
plt.title('Actual vs Predicted Sales')
plt.show()
```

	ID	TV	Radio	Newspaper	Sales
0	1	187.203376	37.934274	31.679612	13.935866
1	2	228.797342	43.573138	5.452866	21.739807
2	3	200.690844	57.247340	26.412179	9.436481
3	4	186.220796	77.553552	36.894668	14.878905
4	5	155.913700	40.015672	6.978894	23.592375

Training Set Shape: (350, 3) (350,)
Testing Set Shape: (150, 3) (150,)
Model Coefficients: [-0.00400006 -0.00810181 0.01809161]
Model Intercept: 15.87427913070574
R-squared: -0.01
Mean Squared Error: 32.80



Question 2)

```
# 1. Import Libraries
import pandas as pd
import numpy as np
import seaborn as sn
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression

# 2. Creating the 'realestate' Dataset
np.random.seed(0) # For reproducibility
realestate_data = pd.DataFrame({
    'ID': range(1, 501),
    'Flat': np.random.uniform(50, 500, 500),
    'Houses': np.random.uniform(20, 300, 500),
    'Purchases': np.random.uniform(5, 50, 500)
})
print(realestate_data.head()) # Previewing the first 5 rows

# 3. Splitting the Data into Independent and Target Variables
X = realestate_data[['Flat', 'Houses']].values # Independent
variables
y = realestate_data['Purchases'].values # Target variable

# Splitting into Training and Testing Sets (7:3 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=0)

print("Training Set Shape:", X_train.shape, y_train.shape)
print("Testing Set Shape:", X_test.shape, y_test.shape)

# 4. Building a Linear Regression Model
linear_regressor = LinearRegression() # Creating the linear
regression model
linear_regressor.fit(X_train, y_train) # Fitting the model

# Making predictions on the test set
y_pred = linear_regressor.predict(X_test)

# 5. Evaluating the Model
print("Model Coefficients:", linear_regressor.coef_)
print("Model Intercept:", linear_regressor.intercept_)

# Calculating R-squared and Mean Squared Error
from sklearn.metrics import r2_score, mean_squared_error
r2 = r2_score(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)

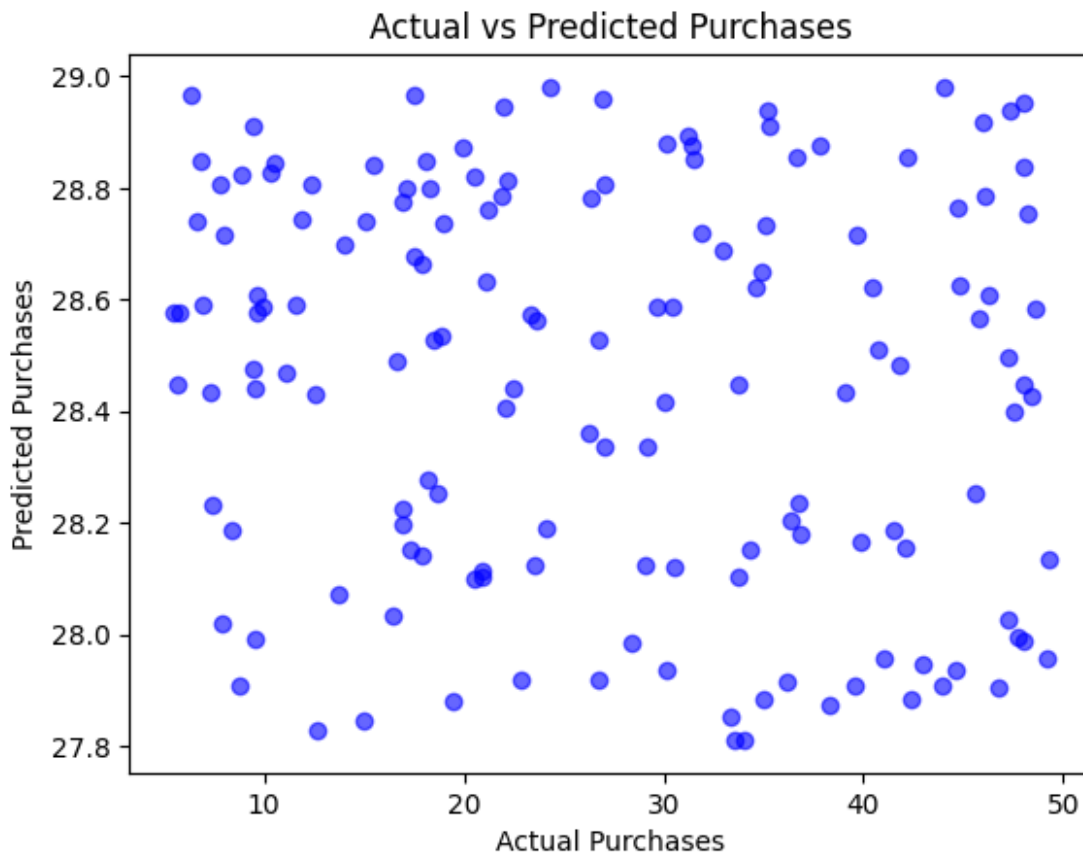
print(f"R-squared: {r2:.2f}")
```

```
print(f"Mean Squared Error: {mse:.2f}")

# Plotting Actual vs Predicted Purchases
plt.scatter(y_test, y_pred, color='blue', alpha=0.6)
plt.xlabel('Actual Purchases')
plt.ylabel('Predicted Purchases')
plt.title('Actual vs Predicted Purchases')
plt.show()
```

	ID	Flat	Houses	Purchases
0	1	296.966077	106.906631	31.679612
1	2	371.835215	124.449762	5.452866
2	3	321.243519	166.991724	26.412179
3	4	295.197432	230.166606	36.894668
4	5	240.644660	113.382090	6.978894

Training Set Shape: (350, 2) (350,)
 Testing Set Shape: (150, 2) (150,)
 Model Coefficients: [0.00023652 -0.0041004]
 Model Intercept: 28.9991454239178
 R-squared: -0.02
 Mean Squared Error: 181.83



Question 3)

```
# 1. Import Libraries
import pandas as pd
import numpy as np
import seaborn as sn
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix

# 2. Creating the 'User' Dataset
np.random.seed(0) # For reproducibility
user_data = pd.DataFrame({
    'User ID': range(1, 501),
    'Gender': np.random.choice(['Male', 'Female'], size=500),
    'Age': np.random.randint(18, 60, size=500),
    'EstimatedSalary': np.random.uniform(20000, 150000, 500),
    'Purchased': np.random.choice([0, 1], size=500, p=[0.7, 0.3])
})
print(user_data.head()) # Previewing the first 5 rows

# Encoding Gender Column (Male: 1, Female: 0)
user_data['Gender'] = user_data['Gender'].map({'Male': 1, 'Female': 0})

# 3. Splitting the Data into Independent and Target Variables
X = user_data[['Gender', 'Age', 'EstimatedSalary']].values # Independent variables
y = user_data['Purchased'].values # Target variable

# Splitting into Training and Testing Sets (7:3 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=0)

print("Training Set Shape:", X_train.shape, y_train.shape)
print("Testing Set Shape:", X_test.shape, y_test.shape)

# 4. Building a Logistic Regression Model
logistic_regressor = LogisticRegression() # Creating the logistic regression model
logistic_regressor.fit(X_train, y_train) # Fitting the model

# Making predictions on the test set
y_pred = logistic_regressor.predict(X_test)

# 5. Evaluating the Model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
```

```

print(f"Accuracy: {accuracy:.2f}")
print("Confusion Matrix:")
print(conf_matrix)

# Visualizing the Confusion Matrix
sn.heatmap(conf_matrix, annot=True, fmt="d", cmap="Blues",
xticklabels=['Not Purchased', 'Purchased'], yticklabels=['Not
Purchased', 'Purchased'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	1	Male	36	93328.366843	0
1	2	Female	41	144401.286761	0
2	3	Female	19	34753.168976	1
3	4	Male	24	55389.816937	0
4	5	Female	48	30667.506897	0

Training Set Shape: (350, 3) (350,)

Testing Set Shape: (150, 3) (150,)

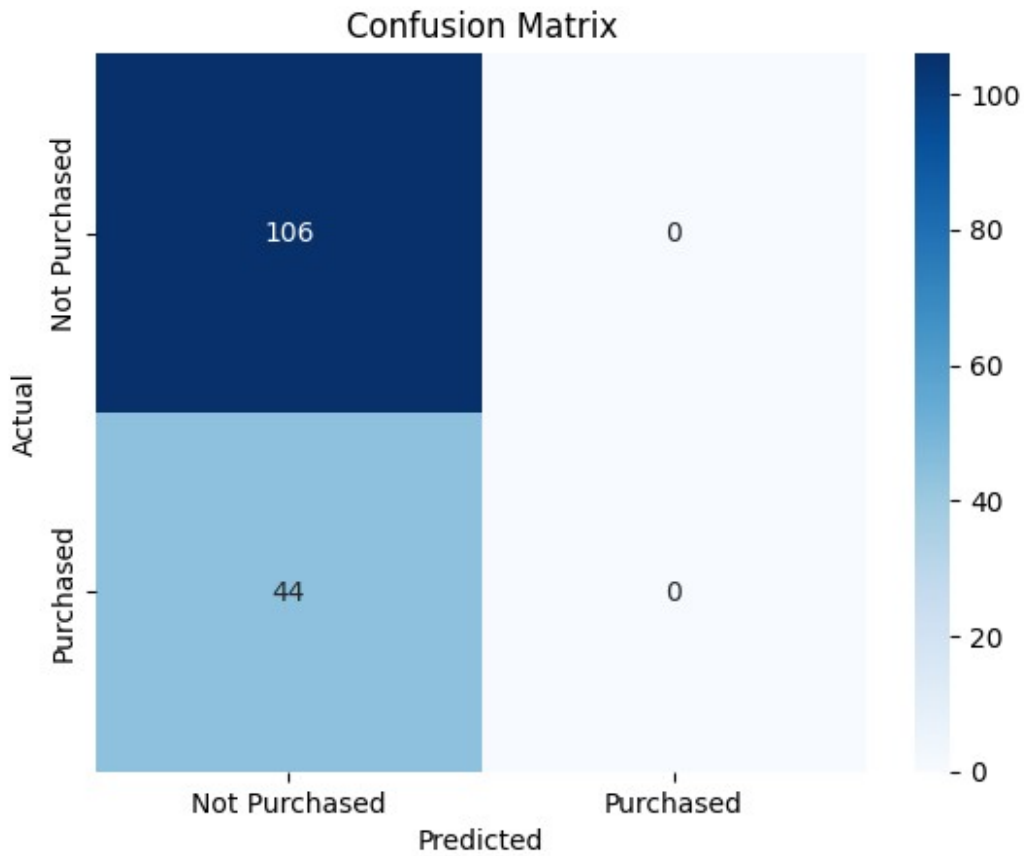
Accuracy: 0.71

Confusion Matrix:

```

[[106  0]
 [ 44  0]]

```



Question 4)

1. Import Libraries

```
import pandas as pd
import numpy as np
import seaborn as sn
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

2. Load the Fish Dataset

Assuming the dataset is downloaded and saved locally as 'Fish.csv'

```
data_path = 'C:\\Users\\ecs\\OneDrive\\Videos\\Documents\\Desktop\\dataset\\Fish.csv'
```

```
data = pd.read_csv(data_path)
```

```
print(data.head()) # Previewing the first 5 rows
```

3. Splitting the Data into Independent and Target Variables

Independent variables: Length1, Length2, Length3, Height, Width

Target variable: Weight

```
X = data[['Length1', 'Length2', 'Length3', 'Height', 'Width']].values
```

```
y = data['Weight'].values
```

```

# Splitting into Training and Testing Sets (7:3 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=0)

print("Training Set Shape:", X_train.shape, y_train.shape)
print("Testing Set Shape:", X_test.shape, y_test.shape)

# 4. Building a Linear Regression Model
linear_regressor = LinearRegression() # Creating the linear
regression model
linear_regressor.fit(X_train, y_train) # Fitting the model

# Making predictions on the test set
y_pred = linear_regressor.predict(X_test)

# 5. Evaluating the Model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse:.2f}")
print(f"R2 Score: {r2:.2f}")

# Visualizing Actual vs Predicted
plt.scatter(y_test, y_pred, color="blue")
plt.xlabel('Actual Weight')
plt.ylabel('Predicted Weight')
plt.title('Actual vs Predicted Weight')
plt.show()

```

	Species	Weight	Length1	Length2	Length3	Height	Width
0	Bream	242.0	23.2	25.4	30.0	11.5200	4.0200
1	Bream	290.0	24.0	26.3	31.2	12.4800	4.3056
2	Bream	340.0	23.9	26.5	31.1	12.3778	4.6961
3	Bream	363.0	26.3	29.0	33.5	12.7300	4.4555
4	Bream	430.0	26.5	29.0	34.0	12.4440	5.1340

Training Set Shape: (111, 5) (111,)

 Testing Set Shape: (48, 5) (48,)

 Mean Squared Error: 32509.60

 R2 Score: 0.81


```

# 1. Import Libraries
import pandas as pd
import numpy as np
import seaborn as sn
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report

# 2. Load the Iris Dataset
from sklearn.datasets import load_iris
data = load_iris()
iris_df = pd.DataFrame(data=data.data, columns=data.feature_names)
iris_df['species'] = data.target

print(iris_df.head()) # Previewing the first 5 rows

# Viewing basic statistical details for each species
for species in range(3):
    print(f"Statistics for {data.target_names[species]}:")
    print(iris_df[iris_df['species'] == species].describe())
    print("\n")

# 3. Splitting the Data into Independent and Target Variables
X = iris_df.iloc[:, :-1].values # Features: sepal and petal lengths
and widths
y = iris_df['species'].values # Target: species

# Splitting into Training and Testing Sets (7:3 ratio)
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.3, random_state=0)

print("Training Set Shape:", X_train.shape, y_train.shape)
print("Testing Set Shape:", X_test.shape, y_test.shape)

# 4. Building a Logistic Regression Model
logistic_regressor = LogisticRegression(max_iter=200) # Creating the
logistic regression model
logistic_regressor.fit(X_train, y_train) # Fitting the model

# Making predictions on the test set
y_pred = logistic_regressor.predict(X_test)

# 5. Evaluating the Model
accuracy = accuracy_score(y_test, y_pred)
conf_matrix = confusion_matrix(y_test, y_pred)
class_report = classification_report(y_test, y_pred,
target_names=data.target_names)

print(f"Accuracy: {accuracy:.2f}")

```

```

print("Confusion Matrix:")
print(conf_matrix)
print("Classification Report:")
print(class_report)

# Visualizing the Confusion Matrix
plt.figure(figsize=(8, 6))
sn.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
xticklabels=data.target_names, yticklabels=data.target_names)
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width
(cm) \				
0	5.1	3.5	1.4	
0.2				
1	4.9	3.0	1.4	
0.2				
2	4.7	3.2	1.3	
0.2				
3	4.6	3.1	1.5	
0.2				
4	5.0	3.6	1.4	
0.2				

	species
0	0
1	0
2	0
3	0
4	0

Statistics for setosa:

	sepal length (cm)	sepal width (cm)	petal length (cm)	\
count	50.00000	50.000000	50.000000	
mean	5.00600	3.428000	1.462000	
std	0.35249	0.379064	0.173664	
min	4.30000	2.300000	1.000000	
25%	4.80000	3.200000	1.400000	
50%	5.00000	3.400000	1.500000	
75%	5.20000	3.675000	1.575000	
max	5.80000	4.400000	1.900000	

	petal width (cm)	species
count	50.000000	50.0
mean	0.246000	0.0
std	0.105386	0.0
min	0.100000	0.0
25%	0.200000	0.0

50%	0.200000	0.0
75%	0.300000	0.0
max	0.600000	0.0

Statistics for versicolor:

	sepal length (cm)	sepal width (cm)	petal length (cm)	\
count	50.000000	50.000000	50.000000	
mean	5.936000	2.770000	4.260000	
std	0.516171	0.313798	0.469911	
min	4.900000	2.000000	3.000000	
25%	5.600000	2.525000	4.000000	
50%	5.900000	2.800000	4.350000	
75%	6.300000	3.000000	4.600000	
max	7.000000	3.400000	5.100000	

	petal width (cm)	species
count	50.000000	50.0
mean	1.326000	1.0
std	0.197753	0.0
min	1.000000	1.0
25%	1.200000	1.0
50%	1.300000	1.0
75%	1.500000	1.0
max	1.800000	1.0

Statistics for virginica:

	sepal length (cm)	sepal width (cm)	petal length (cm)	\
count	50.000000	50.000000	50.000000	
mean	6.588000	2.974000	5.552000	
std	0.63588	0.322497	0.551895	
min	4.900000	2.200000	4.500000	
25%	6.22500	2.800000	5.100000	
50%	6.50000	3.000000	5.550000	
75%	6.90000	3.175000	5.875000	
max	7.90000	3.800000	6.900000	

	petal width (cm)	species
count	50.000000	50.0
mean	2.02600	2.0
std	0.27465	0.0
min	1.40000	2.0
25%	1.80000	2.0
50%	2.00000	2.0
75%	2.30000	2.0
max	2.50000	2.0

Training Set Shape: (105, 4) (105,)

Testing Set Shape: (45, 4) (45,)

Accuracy: 0.98

Confusion Matrix:

```
[[16  0  0]
```

```
 [ 0 17  1]
```

```
 [ 0  0 11]]
```

Classification Report:

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	16
versicolor	1.00	0.94	0.97	18
virginica	0.92	1.00	0.96	11
accuracy			0.98	45
macro avg	0.97	0.98	0.98	45
weighted avg	0.98	0.98	0.98	45

